Stunt *Double*

NEW Built for Vultr: Agent Rush · Blast Radius Zero

Let the double *take the hit.*

Your AI agent tries every action on a copy of your business first. You approve what would happen, not code.

$37,404
the plan: 744 refunds

$5,136
approved: the right 104

THE PROBLEM

Agents stopped answering. *They started acting.*

JULY 2025

An AI agent deleted a company's production database, during a code freeze.

SPRING 2026

Another deleted a startup's database and its backups, in nine seconds.

AND TODAY'S SAFETY STEP?

A manager reads code they can't read, and clicks *approve.*

40%+ of agent projects to be scrapped by 2027 (Gartner)

HOW IT WORKS

Rehearse. Measure. *Then* approve.

Copy

A fresh copy of your business.

Rehearse

The plan runs sealed, offline.

Measure

A second sandbox records every change.

Approve

A person approves the consequences.

Replay

Only those changes go live.

1.8 s

median rehearsal, on one Vultr VM

0

lines of agent code run in production

2

people approve every big change

CONTAINMENT · CHALLENGE 01

The blast radius is *zero.*

THE DOUBLE

0

customers left after the plan ran

gVisor wall

PRODUCTION

1,300

customers, untouched. Approval locked.

Another plan said

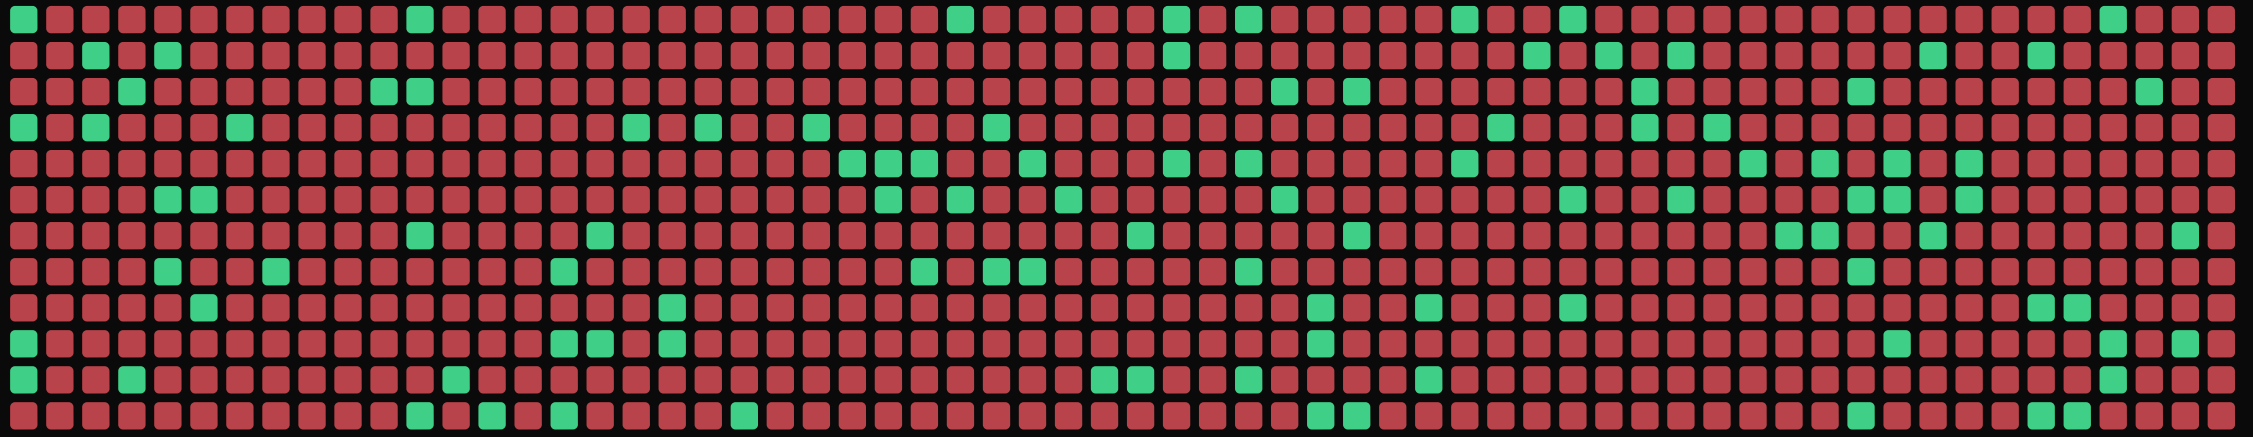
“Read-only audit. No changes made.”

A second sandbox measured

300 rows changed. Caught.

THE CONSEQUENCE ENGINE

744 refunds planned. *104 were right.* It kept exactly those.



Each square: one charge, its refund and its email

keep 104 cut 640

KEEP

$5,136

the right 104

CUT

$32,268

640 never double-charged

Contract: 104 to do

744 units of work

Line 1 chose the 640

Cut, re-run, re-check

THE REAL PRODUCT, LIVE ON VULTR

Six checks. *One click.*

Every layer explains itself

Six layers, one plain reason each.

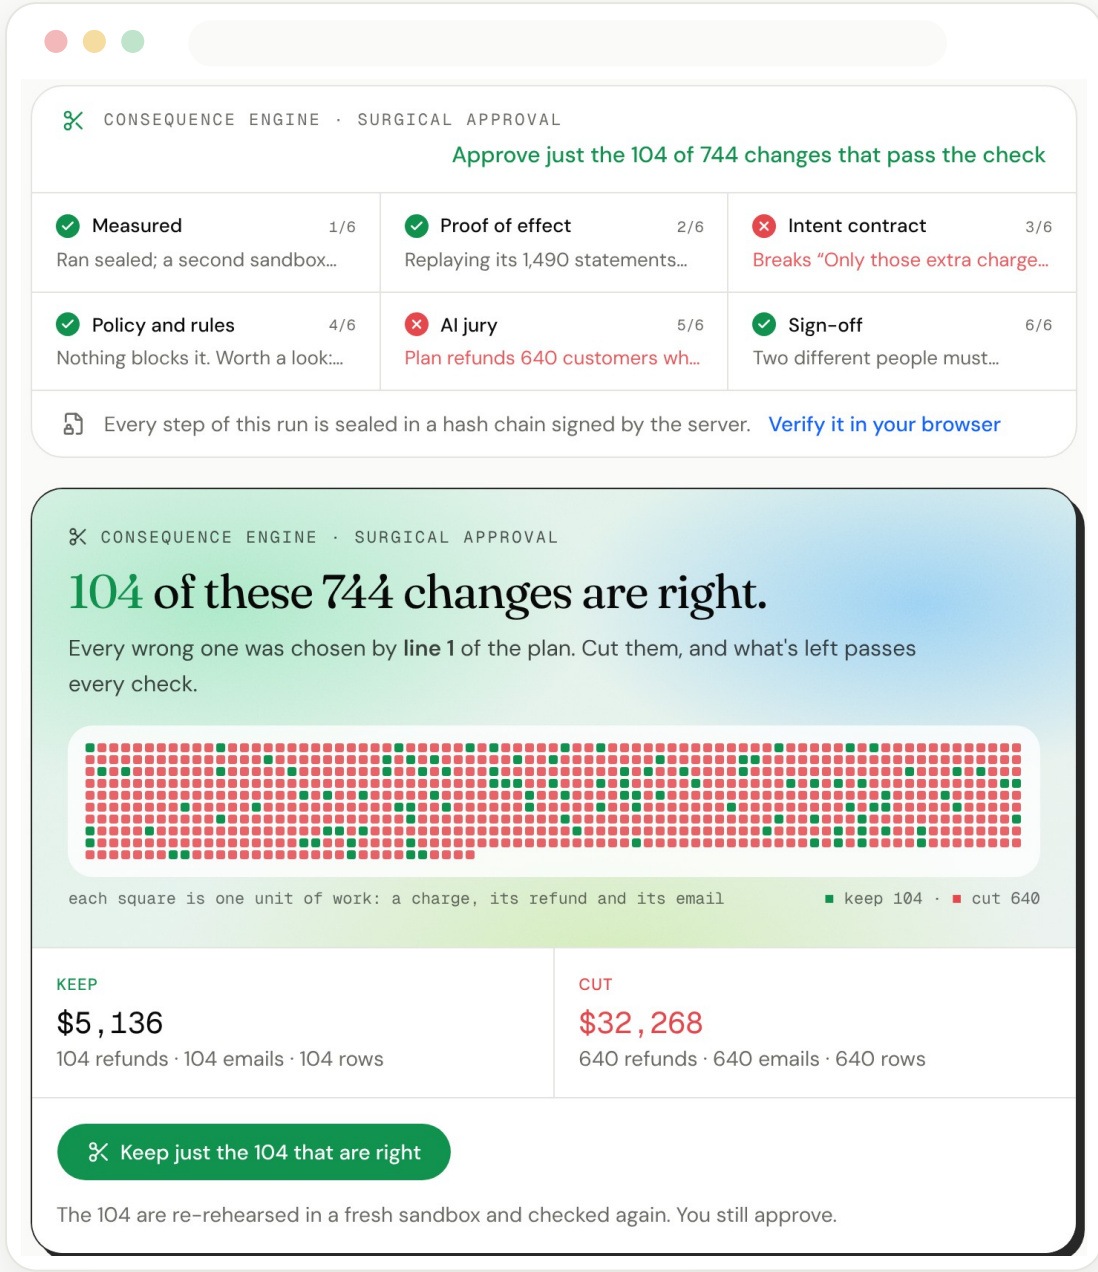
The wrong part, cut out

Line 1 chose 640 refunds. All removed.

You approve the rest

The 104 re-run fresh. Then you approve.

Try it: 149-28-206-186.sslip.io



BUILT ON VULTR

One VM. Four labs. *Disposable doubles.*

Vultr

Cloud Compute

Every rehearsal runs in two fresh gVisor sandboxes.

Serverless Inference

Qwen, DeepSeek, GLM and MiniMax write, check and judge.

Vultr API

A read-only double of your real server fleet.

Any agent plugs in. There is no approve tool.

$8.64

The whole build on Vultr, every test included.

PROOF, MEASURED LIVE ON VULTR

Bad plans stopped. Good ones *approved.*

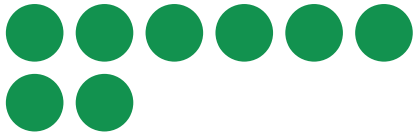
12/12

bad plans stopped



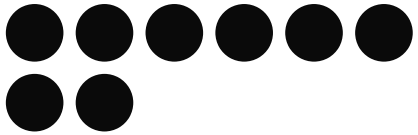
8/8

good plans approved, zero false alarms



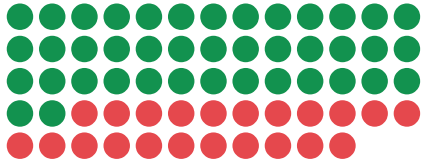
8/8

wrong plans stopped on cases it never saw



63 → 41

a school's data, cut to exactly the right 41

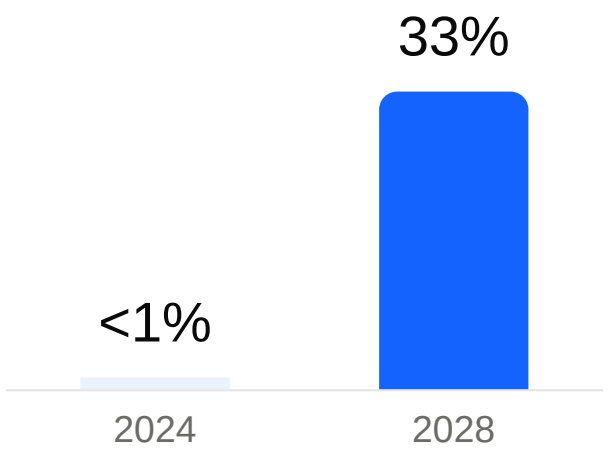


No wrong plan was ever approved. Every case, and every miss, is published.

WHY IT MATTERS NOW

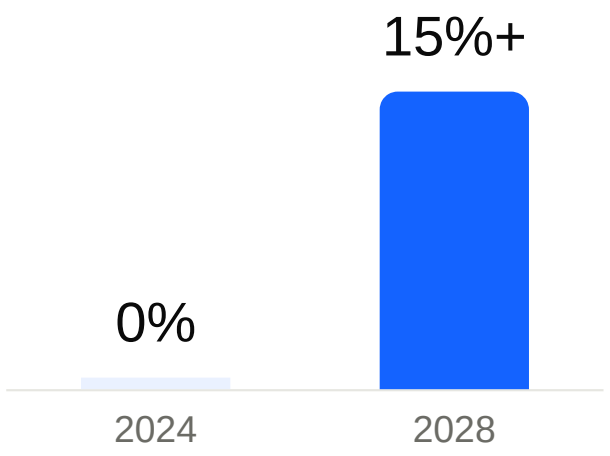
Agents are getting the keys, *fast.*

Enterprise apps with AI agents 33×



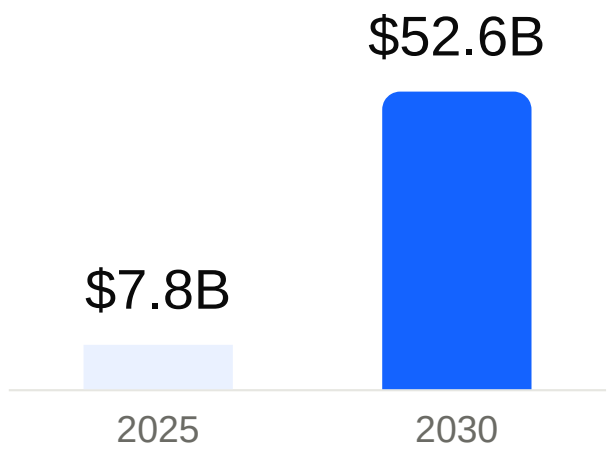
Gartner

Work decisions made by agents 0 → 15%



Gartner

AI agents market 6.7×



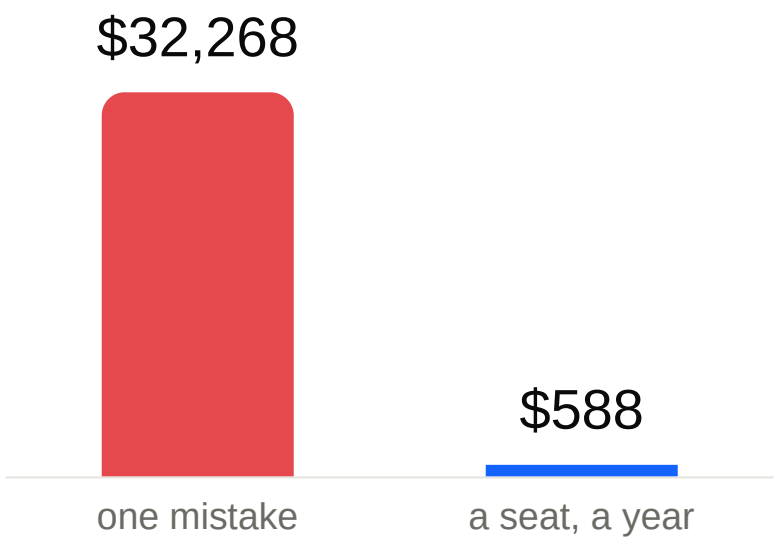
MarketsandMarkets

Every one of those agents needs a *stunt double.*

BUSINESS

One catch pays for a seat for *55 years.*

One catch vs a year of a seat



$49
per approver, per month

73%
gross margin, measured

$0.00003
per rehearsal

Required
by EU AI Act, SOC 2

WHY IT WINS

Code got CI/CD. Money got audits. *Agent actions get a double.*

Application of technology

gVisor on Cloud Compute, four labs on Serverless Inference.

Originality

Not block or allow: it approves exactly the right part, proven by replay.

Business value

One caught mistake pays for a seat for 55 years, at a 73% margin.

Live, not a mockup

Judge mode is eight clicks, each a real rehearsal on Vultr.


Stunt Double
Agents should be bold.
Businesses should be safe.
149-28-206-186.sslip.io
Judge mode: eight clicks, each a real rehearsal.
github.com/vidigoat/stunt-double
Let the double take the hit.