Stunt *Double*

NEW Built for Vultr: Agent Rush · Blast Radius Zero

Let the double *take the hit.*

Your AI agent tries every action on a throwaway copy of your business first. You see what would happen, and approve the **consequences**, not the code.

`$37,404 → 744 refunds caught · $5,136 → the right 104`

THE PROBLEM

Agents stopped answering. *They started acting.*

JULY 2025

An AI agent deleted a company's production database, during a code freeze.

SPRING 2026

Another deleted a startup's database and its backups, in nine seconds.

AND TODAY'S SAFETY STEP?

A manager reads code they can't read, and clicks *approve.*

40%+

of agentic AI projects will be cancelled by 2027, partly for weak risk controls (Gartner)

HOW IT WORKS

Rehearse. Measure. *Then* approve.

Copy

A fresh copy of the data the agent needs.

Rehearse

Sandbox 1 runs the plan. gVisor, no network, no secrets.

Measure

Sandbox 2, our code only, measures what changed.

Approve

A person approves the consequences, in plain English.

Replay

Only those effects reach production, after a drift check.

1.8 s

median rehearsal, both sandboxes, on one $20/month Vultr VM

0

lines of the agent's code ever run against production

2

different people must approve big changes; every undo is rehearsed too

CONTAINMENT · CHALLENGE 01

The blast radius is *zero.*

SANDBOX · CUSTOMERS

0

the plan deleted all 1,300, on the double

PRODUCTION · CUSTOMERS

1,300

untouched. Approval locked.

ANOTHER PLAN SAID

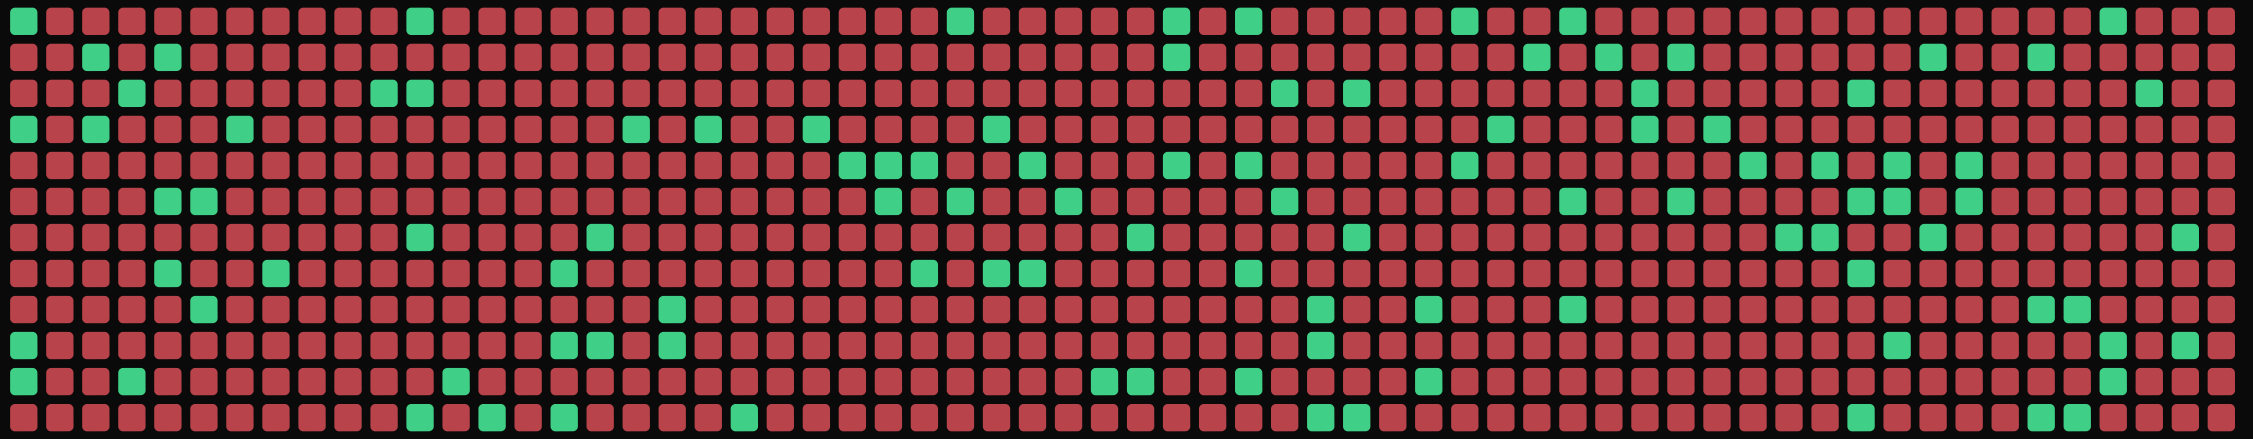
"Audit complete. Read-only: no changes were made."

A SEPARATE SANDBOX MEASURED

300 rows changed. **The plan lied.**

THE CONSEQUENCE ENGINE

744 refunds planned. *104 were right.* It kept exactly those.



Each square: one charge, its refund and its email

keep 104 cut 640

KEEP

$5,136

the right 104 refunds

CUT

$32,268

640 never double-charged

- A contract before the plan: 104 to do
- Every change linked into 744 units
- Replay proves line 1 chose the 640
- Cut them, re-run, check again

PROOF, MEASURED LIVE ON VULTR

Bad plans stopped. Good ones *approved.*

12/12

bad plans stopped by the engine



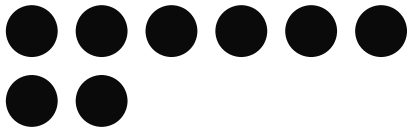
8/8

correct plans approved, 0 false alarms



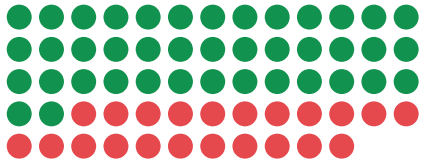
8/8

wrong plans stopped on 16 cases written before the run, never tuned on



63→41

a school's database it had never seen: the trap found, the cut exact

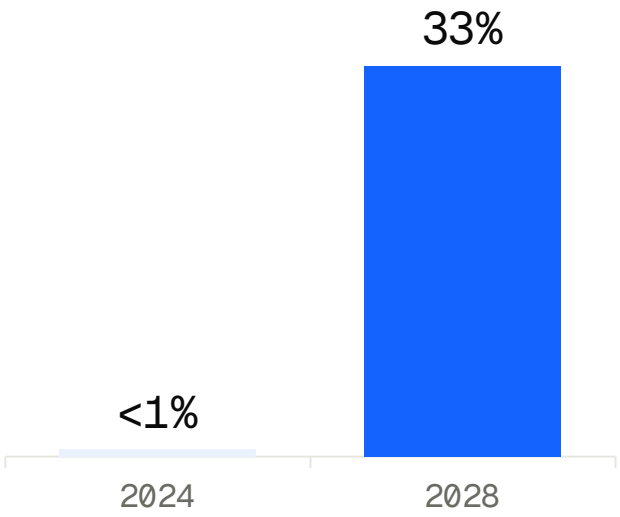


Held-out, honestly: 3 of 8 right plans approved outright; the other 5 were sent to a person with a stated reason. No wrong plan was approved.

WHY IT MATTERS NOW

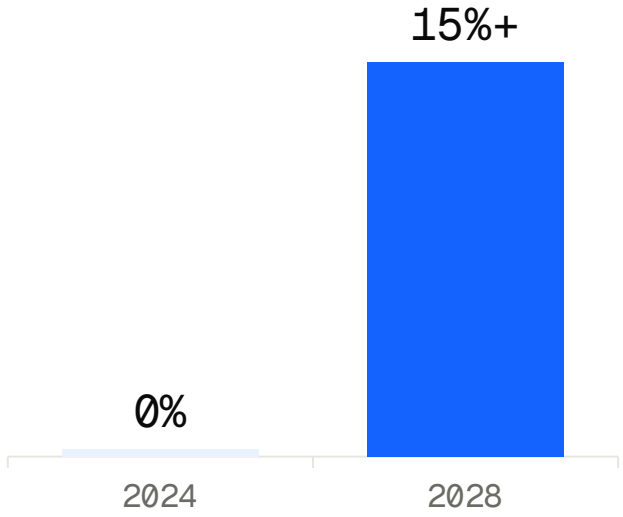
Agents are getting the keys, *fast.*

Enterprise apps with agentic AI



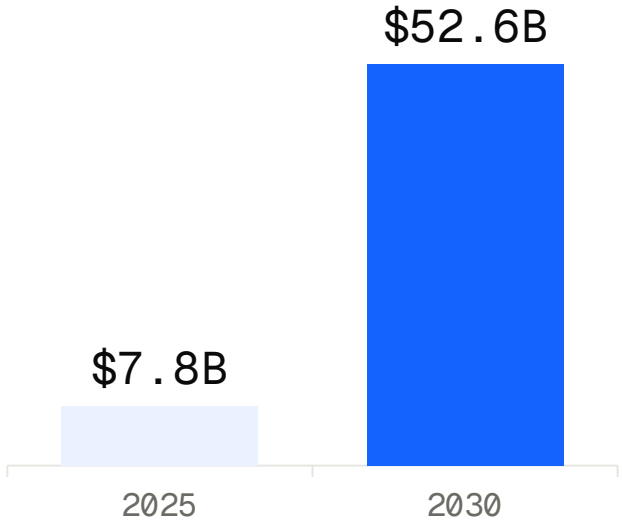
Gartner

Work decisions made by agents



Gartner

AI agents market



MarketsandMarkets

BUILT ON VULTR

One VM. Four labs. *Disposable doubles.*

Vultr

Cloud Compute

Web app, API and MCP server. Every rehearsal runs in two fresh gVisor sandboxes: no network, no secrets.

Serverless Inference

Writes what a request means, cross-checks it, and judges the outcome.

- Qwen · contract
- DeepSeek · cross-check
- GLM · jury + vision
- DeepSeek · jury
- MiniMax · jury

Vultr API

A read-only double of your real server fleet. The key is used once and never kept.

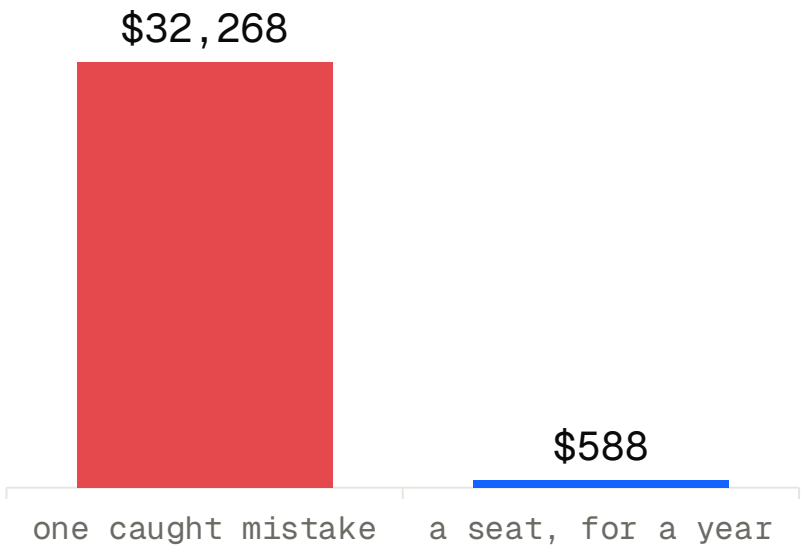
Any agent plugs in: MCP, SDK, CLI or REST. No approve tool.

$8.64 The whole build on Vultr, every test run included.

BUSINESS

One caught mistake pays for a seat for *55 years.*

One catch vs a year of a seat



Unit economics, measured on Vultr

$49
per approver, per month

73%
gross margin at full use

$0.00003
per rehearsal

≤ $0.044
per three-model AI check

Who buys: Heads of Operations, Support and Finance Ops

Why they must: EU AI Act Art. 14 and SOC 2 ask for this record

How it spreads: Open core, MIT: self-host on your Vultr in one command

THE BIGGER PICTURE

Code got CI/CD. Money got audits. *Agent actions get a double.*

TODAY

Live on Vultr

Databases, payments, email, APIs, a real browser and your Vultr fleet.

NEXT

Every system agents touch

More connectors, a Vultr VM per risky rehearsal, approvals in Slack.

THE STANDARD

The approval layer for AI

Every agent action rehearsed, approved, and sealed as evidence.

WHY IT WINS

Technology

gVisor on Cloud Compute, 4 labs on Serverless Inference

Originality

It approves the right part of a plan, proven by replay

Business value

One catch pays a seat for 55 years; 73% margin

Presentation

Live today: judge mode is eight clicks


Stunt Double
Agents should be bold.
Businesses should be safe.
149-28-206-186.sslip.io
Judge mode: eight clicks, each a real rehearsal on Vultr.
github.com/vidigoat/stunt-double
Let the double take the hit.